

Summary

Analytics Engineer with 3+ years of experience designing and scaling enterprise data solutions across the modern data stack. Expert in building high-performance semantic models, ETL/ELT pipelines, and self-service BI architectures using Python, advanced SQL, dbt, and Power BI. Proven ability to translate complex business requirements into optimized dimensional data models and automated workflows that drive measurable impact for stakeholders. Strong track record implementing analytics infrastructure that balances scalability, performance, and maintainability.

Education

Drexel University <i>Master of Science in Business Analytics</i> (GPA: 3.95/4.00)	Philadelphia, PA <i>March 2023</i>
University of Pune <i>Bachelor of Engineering in Electronics & Telecommunications</i> (GPA: 3.75/4.00)	Pune, India <i>June 2020</i>

Experience

HOLMAN Analytics Engineer - Designed enterprise semantic model (star schema) consolidating disparate sources, serving 400+ users and enabling self-service BI - Optimized dbt transformation layer on Databricks using incremental models and materialized views, reducing ETL costs by 20% - Built reusable Power BI semantic models across business domains with DAX calculation groups and query optimization, improving report load times by 40% - Implemented CI/CD pipeline for Power BI (TMDL/PBIP, Git), reducing manual publishing errors by 90% and establishing deployment automation - Led numerous requirements gathering sessions with business stakeholders, translating complex analytical needs into scalable semantic model architecture - Pioneered dbt adoption and best practices across the team, establishing documentation standards and accelerating team capability development	Mt Laurel, New Jersey <i>July 2023 – Present</i>
COLLABERA Data Analyst Intern - Developed predictive churn model (Python: Scikit-learn, Pandas) achieving 78% accuracy to identify at-risk customers for retention campaigns - Consolidated 8+ data sources via SQL and Power BI, optimizing queries and reducing dashboard refresh time by 30% - Built reusable DAX calculation groups for time-intelligence logic, standardizing formulas across Power BI reports	Basking Ridge, New Jersey <i>May 2023 – July 2023</i>
LABWARE Data Science Intern - Built ETL pipeline processing 2M+ crime records using Python (Pandas), AWS S3, and Snowflake for centralized analytics - Trained XGBoost classification model (80% accuracy) to identify crime types and high-risk geographic patterns - Created Tableau dashboards visualizing ML predictions and crime hotspots to support law enforcement decision-making	Wilmington, Delaware <i>June 2022 – September 2022</i>

Projects

H-1B Wage Map - Built Python ETL pipeline processing 200K+ historical H-1B filings into static JSON shards, achieving O(1) lookup for real-time wage intelligence queries - Engineered interactive geospatial map using Mapbox GL JS with client-side GeoJSON rendering for sub-second county-level drill-down across 3,200+ US regions - Integrated OpenAI API with RAG for AI-powered SOC code discovery matching job descriptions to official O*NET occupations
Personal AI Powered Portfolio - Engineered recruiter-focused job-fit analyzer powered by OpenAI Responses API to parse job descriptions and match against structured portfolio evidence - Developed AkshayGPT chat assistant with centralized portfolio knowledge retrieval enabling conversational exploration of projects and technical competencies - Architected responsive Next.js 15 (React 19, TypeScript, App Router, Tailwind CSS) platform with dual API routes for integrated chat and recruiter workflows

Technical Skills

Core Languages: Python, SQL (Query Optimization, Window Functions, Execution Plans, Complex Joins)
Analytics Platforms & BI: Power BI (Advanced: DAX, Calculation Groups, Performance Analyzer, DAX Studio, TMDL, PBIP, CI/CD), Tableau
Data Engineering & Transformation: dbt, Databricks, Apache Spark (PySpark), ETL/ELT, Incremental Models, Materialized Views, Data Lineage
Cloud & Databases: AWS (S3, Redshift), Microsoft Fabric, Snowflake, PostgreSQL, MS SQL Server, MySQL, Git/GitHub
Data Science & Predictive Modeling: Scikit-learn, Pandas, Predictive Modeling (Classification, Time Series), Generative AI Integration